

# Feature Separation and Recalibration for Adversarial Robustness

CVPR 2023 Highlight

[\[paper\]](#) [\[code\]](#)


Woo Jae Kim, Yoonki Cho, Junsik Jung, Sung-Eui Yoon (KAIST)

2023.09.15

Mijin Koo

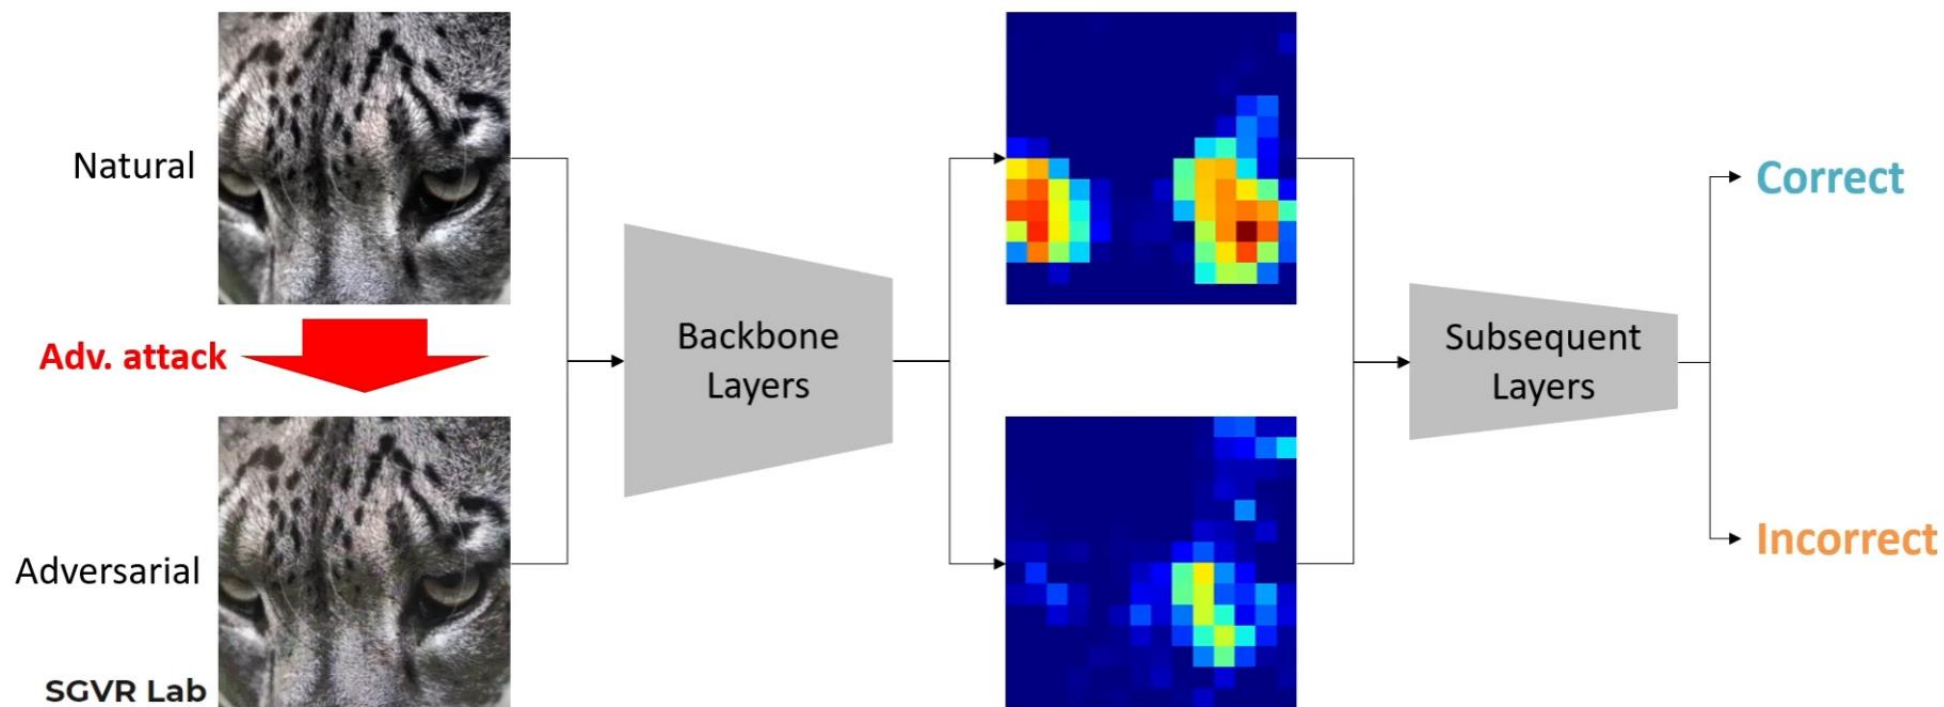
# Summary

- **Feature Activation Disruption upon Adversarial Attack**
  - Feature-level disruptions lead to model mispredictions



# Summary

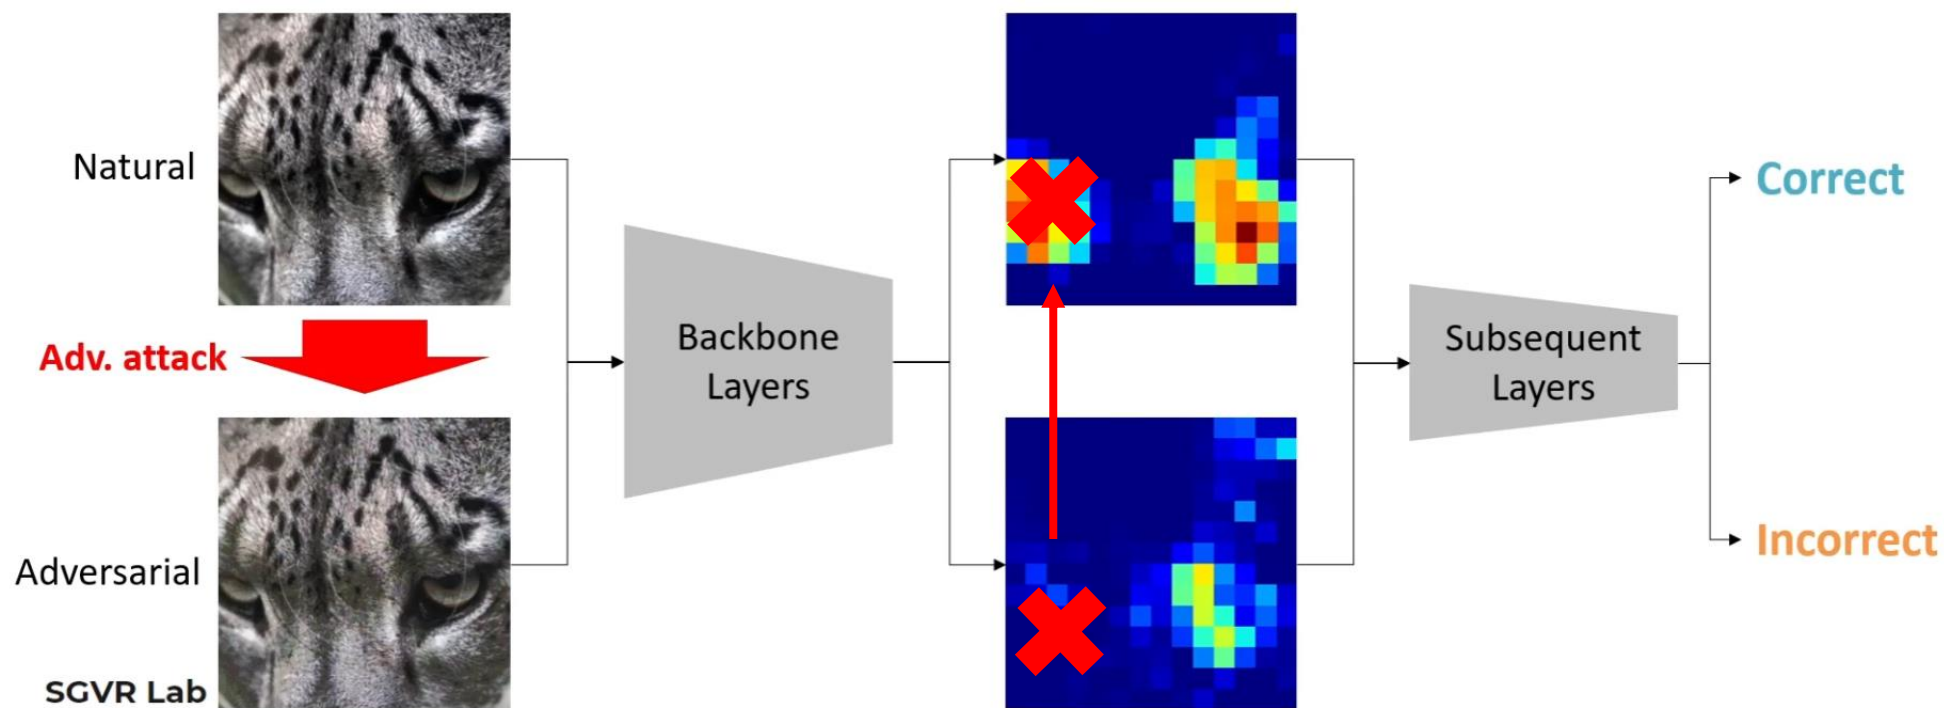
- **Feature Activation Disruption upon Adversarial Attack**
  - Feature-level disruptions lead to model mispredictions



# Summary

## ▪ Limitations of Conventional Defense

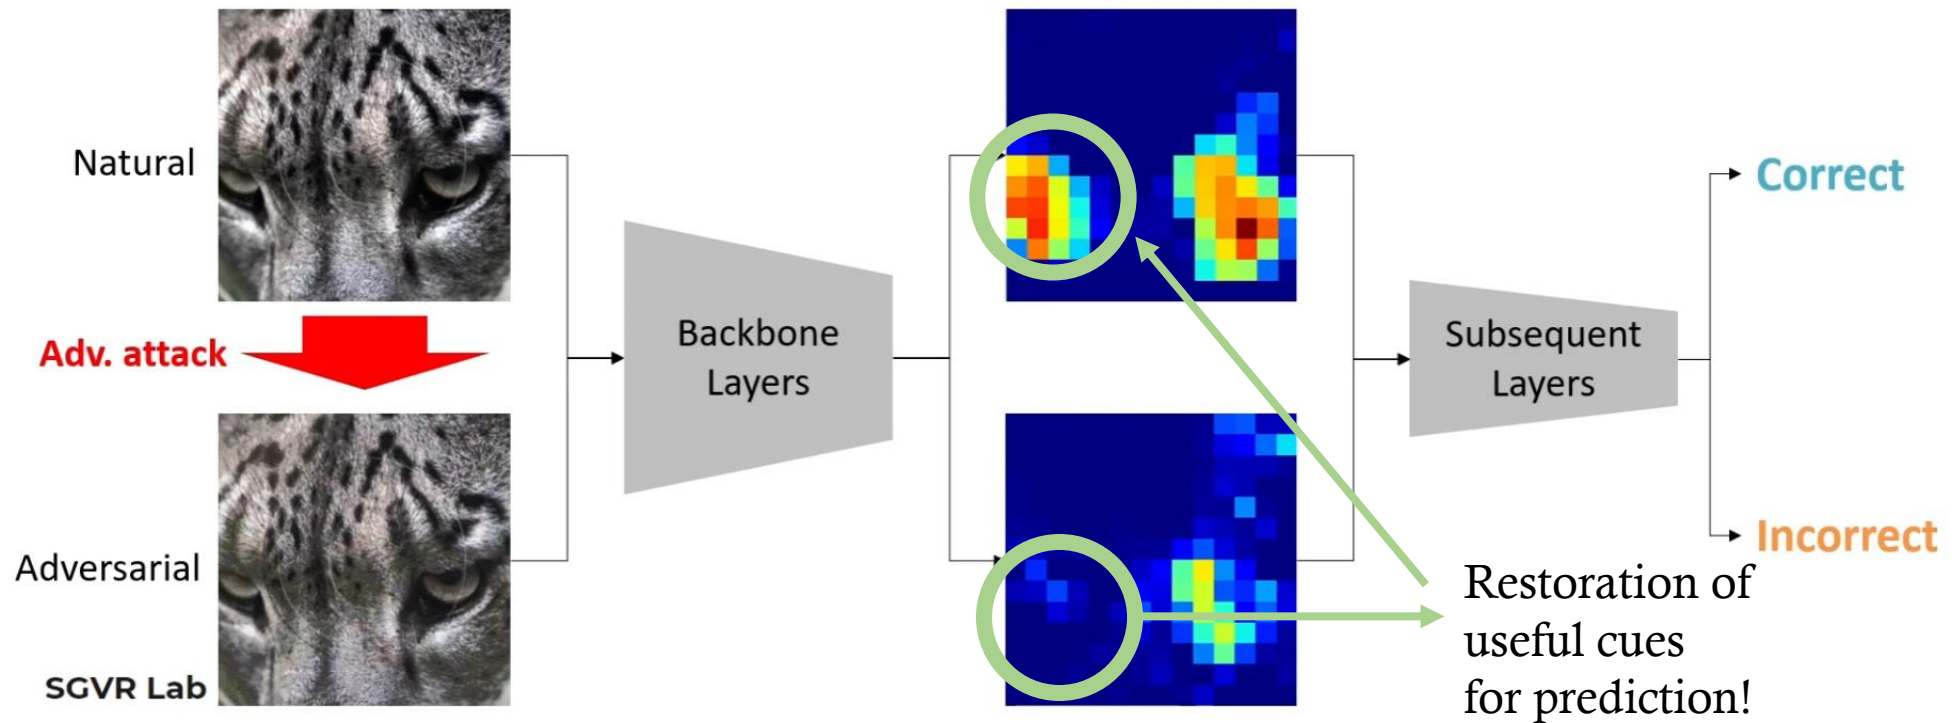
- Conventional defense methods *suppressed or deactivated* disrupted activations
- This approach lead to *loss of potentially discriminative cues*



# Summary

## ■ Proposed Approach

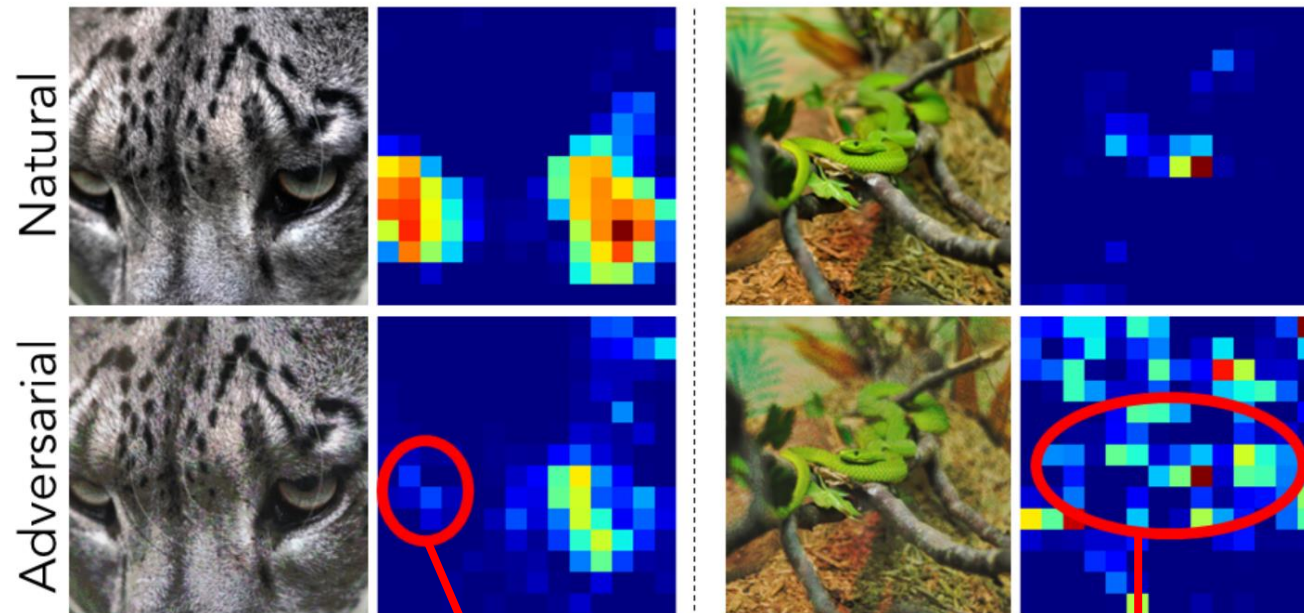
- Instead, we propose to *restore useful cues* from these disrupted activations
- This additional useful cues *enrich* model's ability to make *correct predictions*



# Introduction

## ■ Feature Activation Disruption upon Adversarial Attack

- Adversarial attacks corrupt activations of feature maps
- **Robust feature** with activations that help the model make correct predictions
- **Non-robust feature** with activations that are responsible for model mispredictions upon adversarial attack.



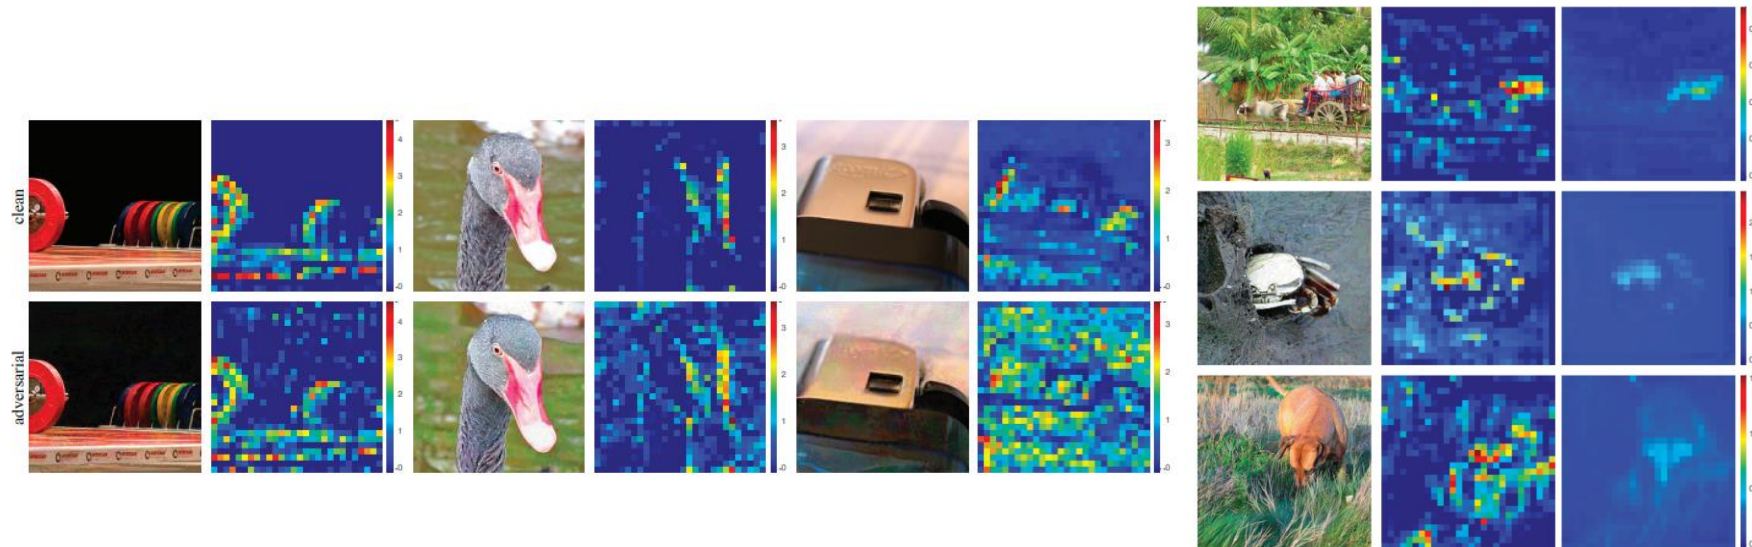
Loss of useful cues  
→ Restoration

Gain of useless information  
→ Deactivating



# Introduction

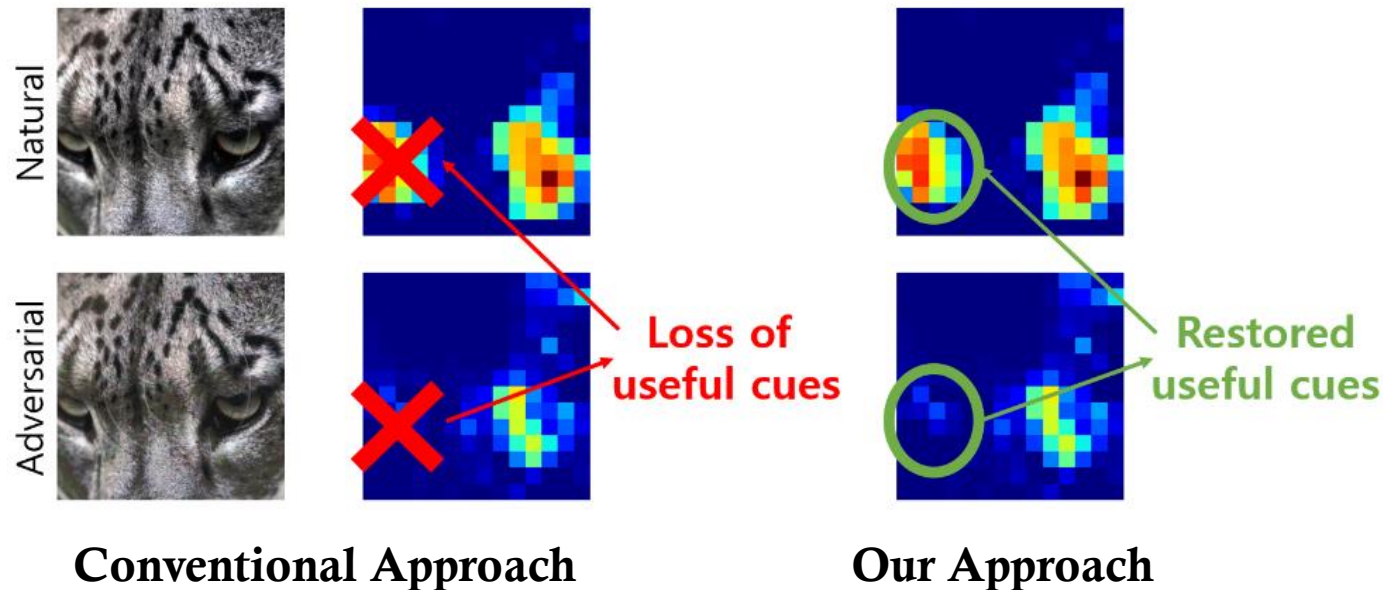
- **Deactivating the non-robust feature activations (Previous)**
  - Adversarial perturbations on images lead to noise in the features.
  - Previous methods solve this problem by deactivating the non-robust feature activations that cause model mispredictions.
  - Increase adversarial robustness by performing *feature denoising*
  - Our networks contain blocks that denoise the features using *non-local means*



Feature denoising for improving adversarial robustness [2019 CVPR]

# Introduction

- **Motivation: non robust feature에도 discriminative cue가 있다!**
  - We propose to *restore useful cues from these disrupted activations* that are otherwise neglected.
- **Contributions**
  - Novel approach of recalibrating deactivated activations to capture useful cues for correct model predictions
  - Easy to plug in Feature and Recalibration(FSR) module
  - Small overhead, successful experiments results

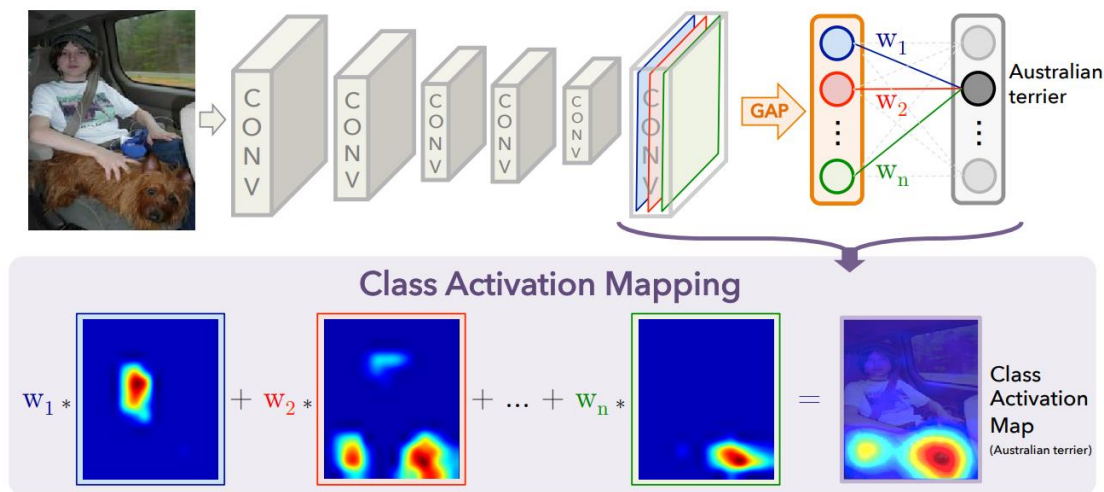




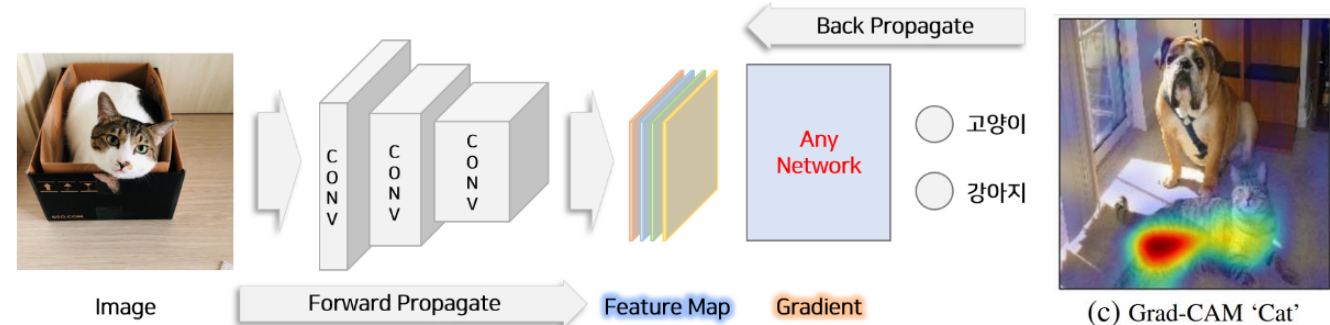
# Introduction

## ■ Visualize Explanation from Deep Networks

- **CAM** (Class Activation Mapping)
  - the predicted class score is mapped back to the previous convolutional layer to generate the class activation maps
  - CAM highlights the class-specific discriminative regions
- **Grad-CAM** (Gradient weighted CAM)
  - uses the gradients of any target flowing into the final convolutional layer



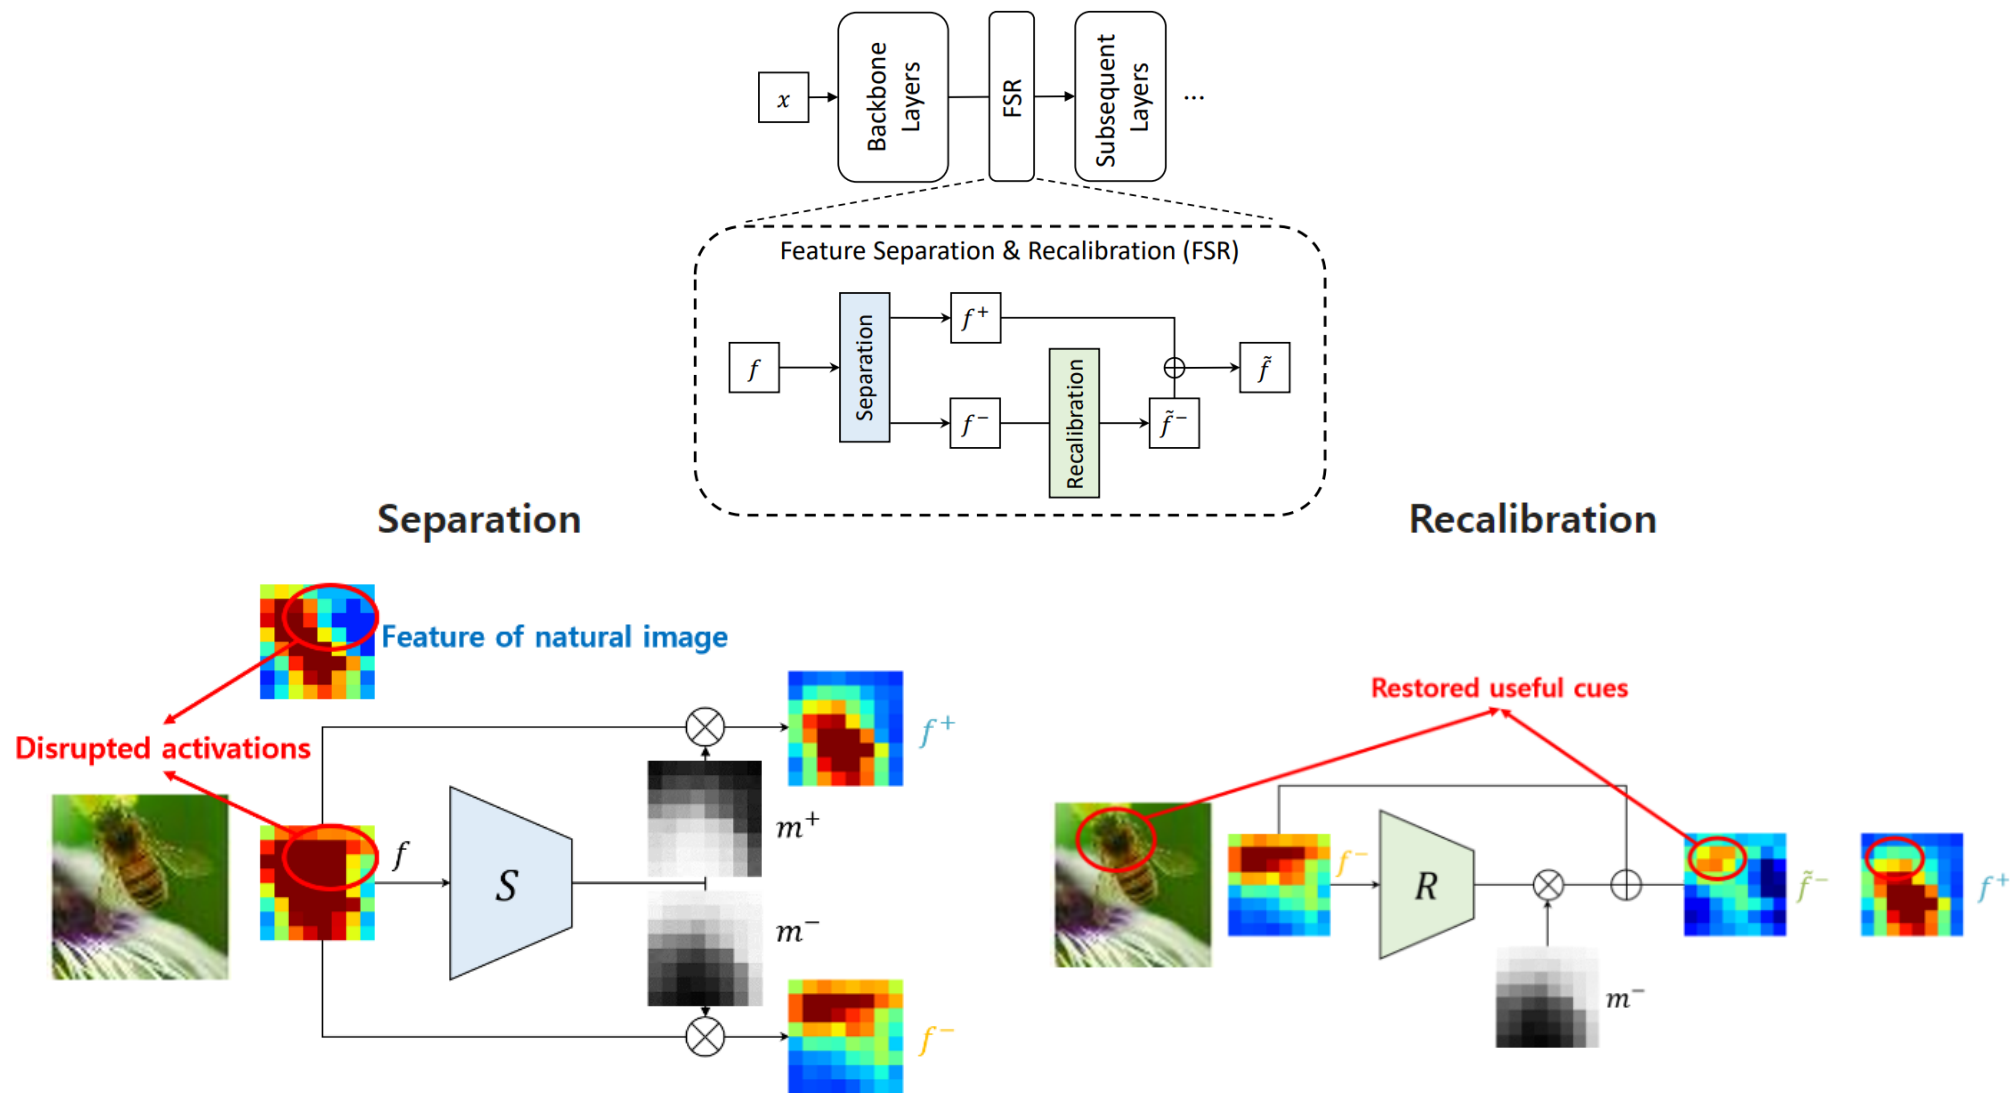
Learning Deep Features for Discriminative Localization [2016 CVPR]



Grad CAM: Visual Explanations from Deep Networks via Gradient-based Localization [2017 ICCV]

# Proposed Approach

- Feature Separation and Recalibration (FSR)



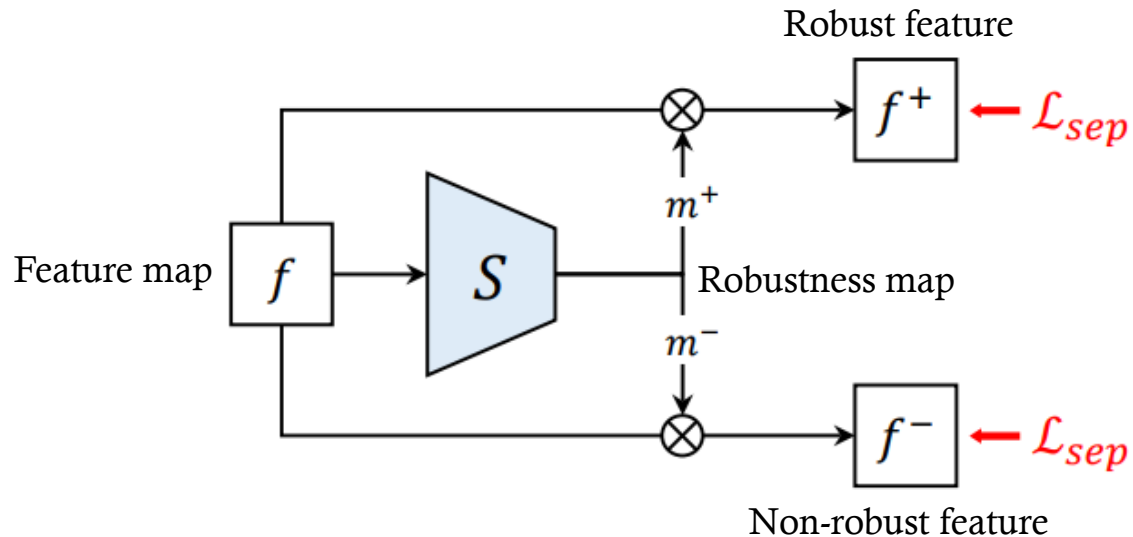
# Proposed Approach

## ■ Separation stage

- Disentangle the feature map into the robust and non-robust features by masking out

## ■ Feature Separation

- Separation Net  $S$
- Input: feature map  $f$
- Output: robustness map  $r$
- Differentiable soft mask  $m \in [0, 1]$ ; approximated by a binary mask  $b \in \{0, 1\}$ 
  - $m^+ = 1 - m^-$



- Robust feature  $f^+ = m^+ \otimes f$
- Non-robust feature  $f^- = m^- \otimes f$

# Proposed Approach

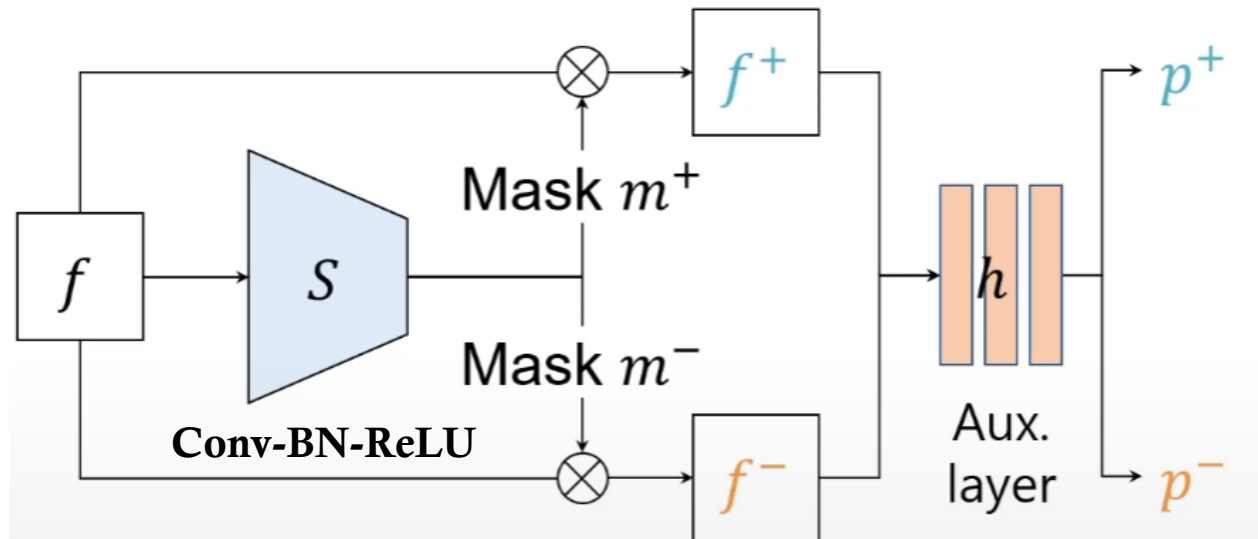
## ■ Feature Separation

- Separation Net  $S$  learns robustness score
- $L_{sep}$  guides the Separation Net to assign high robustness scores to units that help the auxiliary layer make correct predictions

$$\mathcal{L}_{sep} = - \sum_{i=1}^N (y_i \cdot \log(p_i^+) + y'_i \cdot \log(p_i^-)), \quad (3)$$

$$\mathcal{L}_{sep} = \underbrace{\mathcal{H}(p^+, y)}_{\text{Cross-entropy loss}} + \underbrace{\mathcal{H}(p^-, y')}_{\text{Cross-entropy loss}}$$

GT label      Pred. logit      Wrong label

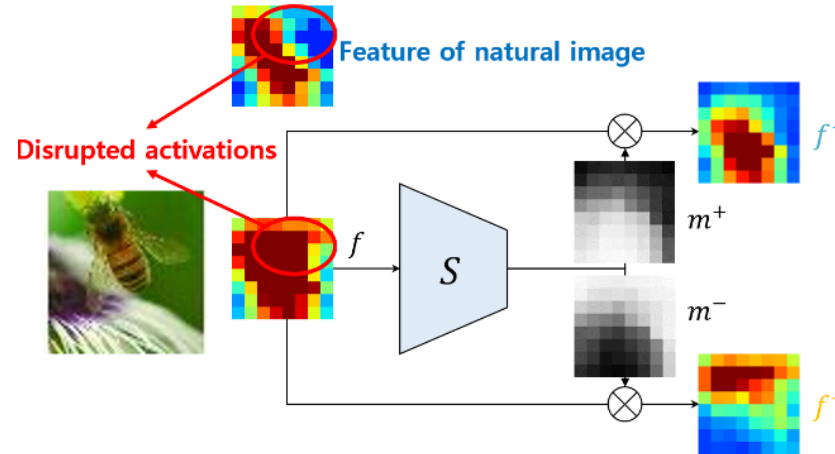


# Proposed Approach

## ■ Feature Separation

### • Soft mask M

- Positive mask emphasizes activations relevant to *correct predictions*
- Negative mask emphasizes activations relevant to *mispredictions*



- $b \in \{0, 1\}$  with a differentiable soft mask  $m \in [0, 1]$ 
  - By Gumbel softmax Approximate a binary mask
  - $r$ : robustness map
  - $g_1, g_2$ : samples from Gumbel distribution such that  $g = -\log(-\log(u))$ ,  $u \sim \text{Uniform}(0, 1)$

$$m = \frac{e^{((\log(\sigma(r)) + g_1)/\tau)}}{e^{((\log(\sigma(r)) + g_1)/\tau)} + e^{((\log(1 - \sigma(r)) + g_2)/\tau)}}, \quad (2)$$

# Proposed Approach

## ■ Feature Separation

### • Gumbel Softmax

- 이산 확률 분포에서의 샘플링을 연속적으로 다루기 위한 방법 → gradient 계산 가능

$$y_i = \frac{\exp((\log(\pi_i) + g_i)/\tau)}{\sum_{j=1}^k \exp((\log(\pi_j) + g_j)/\tau)} \quad p_i = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}}$$

Gumbel softmax

Softmax

- $r$ : robustness map
- $\sigma$ : sigmoid function
- $g$ : samples from Gumbel distribution such that  $g = -\log(-\log(u))$ ,  $u \sim \text{Uniform}(0, 1)$
- $\tau$ : temperature that controls the effect of  $g$

$$m = \frac{e^{((\log(\sigma(r)) + g_1)/\tau)}}{e^{((\log(\sigma(r)) + g_1)/\tau)} + e^{((\log(1 - \sigma(r)) + g_2)/\tau)}}, \quad (2)$$



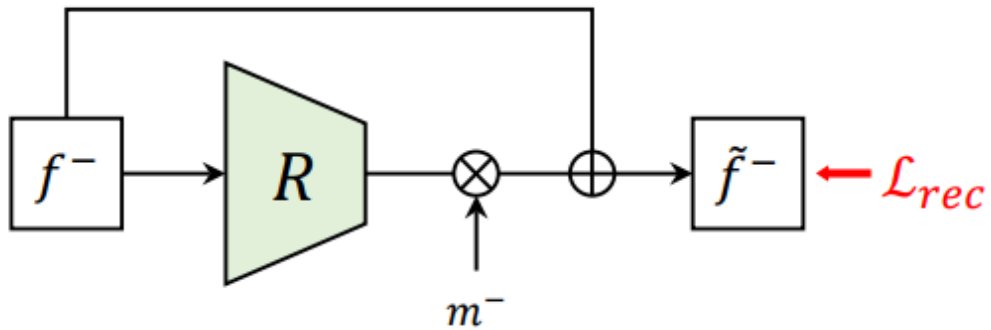
# Proposed Approach

## ■ Recalibration Stage

- Adjust the non-robust feature activations to capture the additional useful cues

## ■ Feature Recalibration

- Recalibration Net  $R$
- Input: non-robust feature map  $f^-$
- Output: recalibrated feature,  $\tilde{f}^- = f^- + m^- \otimes R(f^-)$



# Proposed Approach



## ■ Feature Recalibration

- Guided network  $R$  to restore useful cues relevant to correct prediction
- Input: non-robust feature map  $f^-$
- Output: recalibrated feature,  $\tilde{f}^- = f^- + m^- \otimes R(f^-)$

$$\mathcal{L}_{rec} = - \sum_{i=1}^N y_i \cdot \log(\tilde{p}_i^-), \quad (4)$$

$$\mathcal{L}_{rec} = \mathcal{H}(\tilde{p}^-, y)$$

Cross-entropy      Pred.      GT



# Proposed Approach

## ■ Model Training

- Can be attached to any adversarial training(AT) technique with objective  $L_{cls}$
- FSR is highly modularized and easy to plug-in
- Trained in an end-to-end manner

$$\mathcal{L} = \mathcal{L}_{cls} + \frac{1}{|L|} \sum_{l \in L} (\lambda_{sep} \cdot \mathcal{L}_{sep}^l + \lambda_{rec} \cdot \mathcal{L}_{rec}^l), \quad (5)$$

- $L_{cls}$ : classification loss for adversarial training
- $L_{sep}$ : feature separation loss
- $L_{rec}$ : feature recalibration loss
- $\lambda_{sep}, \lambda_{rec}$ : hyperparameters that control weights

# Experiments

- **Experimental Setups**

- **Baselines**

- PGD adversarial training (AT) [1]
    - TRADES [2]
    - MART [3]

- **Datasets**

- CIFAR-10/100
    - SVHN
    - Tiny ImageNet

- **Models**

- ResNet18
    - VGG16
    - WideResNet-34-10

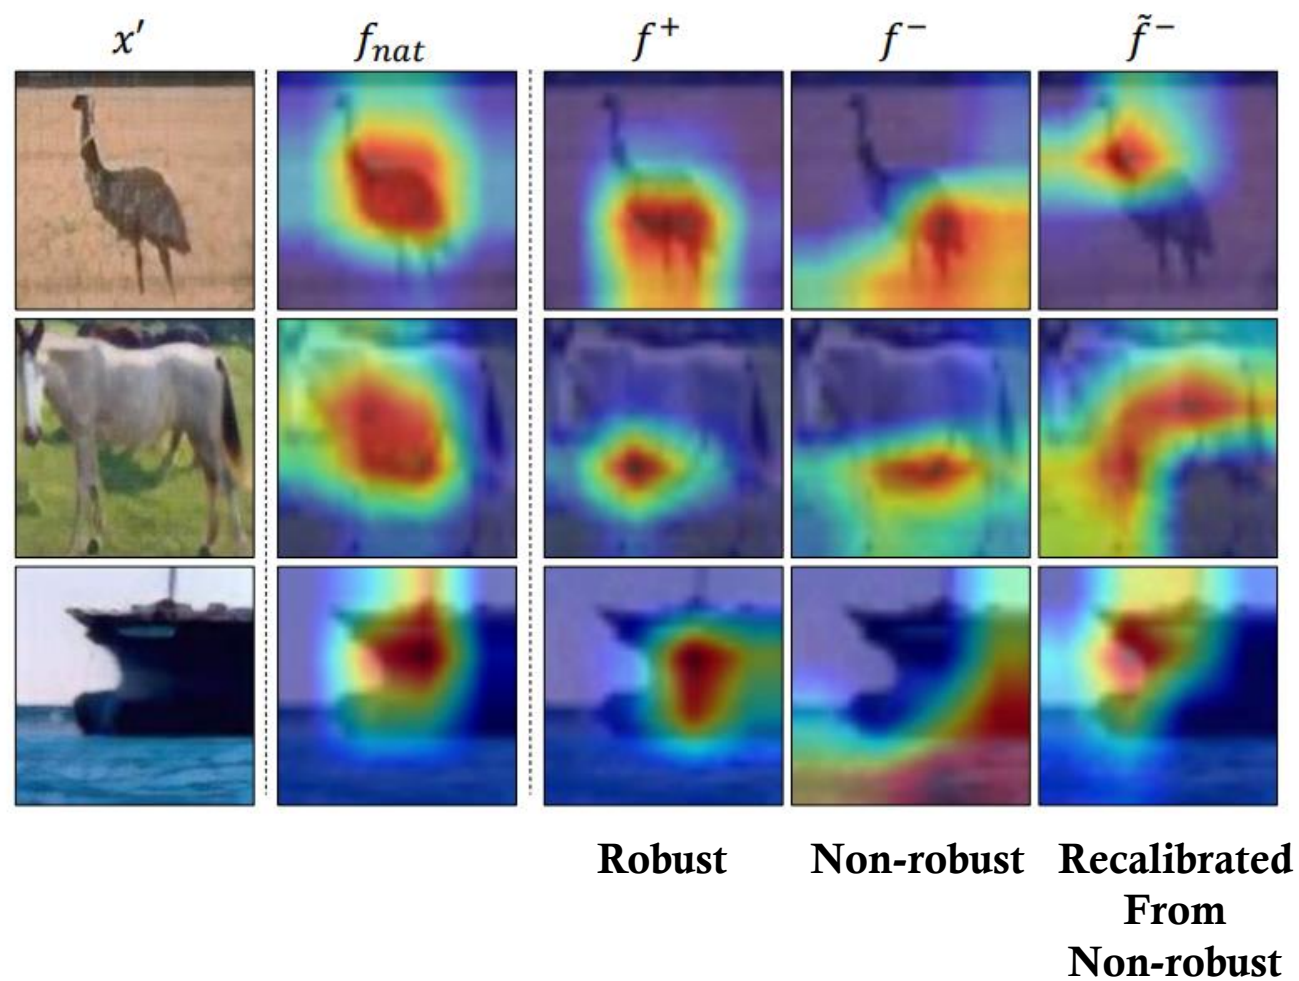
[1] Madry et al., Towards deep learning models resistant to adversarial attacks. [ICLR 2018]

[2] Zhang et al., Theoretically principled trade-off between robustness and accuracy. [ICML 2019]

[3] Wang et al., Improving adversarial robustness via channel-wise activation suppressing. [ICLR 2021]

# Experiments

- Qualitative Results



# Experiments

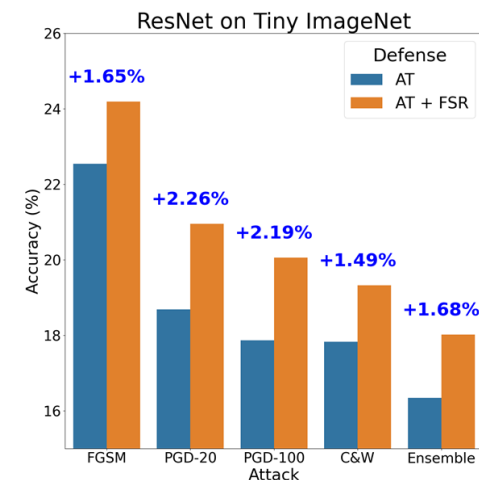

## Application to Adversarial Training

<i>ResNet-18</i>	CIFAR-10						SVHN					
Method	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble
AT	<b>85.02</b>	56.21	48.22	46.37	47.38	45.51	91.21	55.55	40.85	37.54	40.61	37.41
AT + FSR	81.46	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	<b>49.44</b>	<b>48.34</b>	<b>91.28</b>	<b>60.46</b>	<b>43.94</b>	<b>39.01</b>	<b>43.22</b>	<b>38.81</b>
TRADES	<b>86.31</b>	57.21	50.74	49.44	48.66	47.89	90.99	61.31	47.12	43.55	45.48	42.99
TRADES + FSR	84.49	<b>58.29</b>	<b>52.27</b>	<b>51.28</b>	<b>49.92</b>	<b>49.28</b>	<b>91.39</b>	<b>68.85</b>	<b>51.49</b>	<b>47.50</b>	<b>46.70</b>	<b>46.17</b>
MART	82.73	56.65	50.88	49.15	47.21	45.98	<b>90.50</b>	58.21	43.61	40.43	42.20	40.07
MART + FSR	<b>83.28</b>	<b>59.55</b>	<b>54.80</b>	<b>53.69</b>	<b>48.98</b>	<b>48.36</b>	89.87	<b>61.06</b>	<b>46.51</b>	<b>42.94</b>	<b>43.89</b>	<b>42.40</b>

Table 1. Robustness (accuracy (%)) of adversarial training strategies (AT, TRADES, MART) with (+ FSR) and without our FSR module against diverse white-box attacks on ResNet-18. Better results are marked in **bold**.

<i>VGG16</i>	CIFAR-10						SVHN					
Method	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble
AT	<b>80.56</b>	53.47	47.17	45.58	45.82	43.71	89.59	54.88	40.27	36.90	39.46	36.62
AT + FSR	80.06	<b>54.40</b>	<b>49.82</b>	<b>48.82</b>	<b>47.28</b>	<b>46.24</b>	<b>91.44</b>	<b>65.01</b>	<b>45.99</b>	<b>39.07</b>	<b>43.08</b>	<b>38.15</b>
TRADES	<b>82.44</b>	53.92	47.39	46.20	44.80	44.20	90.48	61.50	45.99	40.00	42.82	39.27
TRADES + FSR	80.78	<b>55.48</b>	<b>49.95</b>	<b>49.03</b>	<b>46.28</b>	<b>45.90</b>	<b>91.89</b>	<b>69.25</b>	<b>54.56</b>	<b>47.81</b>	<b>46.66</b>	<b>44.10</b>
MART	76.11	54.86	51.06	50.16	43.53	43.01	89.95	59.03	42.89	38.73	39.12	37.64
MART + FSR	<b>79.18</b>	<b>56.41</b>	<b>52.69</b>	<b>52.13</b>	<b>44.49</b>	<b>44.20</b>	<b>90.60</b>	<b>62.28</b>	<b>47.17</b>	<b>42.50</b>	<b>43.44</b>	<b>40.73</b>

Table 2. Robustness (accuracy (%)) of adversarial training strategies (AT, TRADES, MART) with (+ FSR) and without our FSR module against diverse white-box attacks on VGG16. Better results are marked in **bold**.





# Experiments

## ▪ Comparison with only FSR (w/o Adversarial Training)

Feature deactivation or suppression	Method	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble	AutoAttack
	AT [1]	85.02	56.21	48.22	46.37	47.38	45.51	44.11
	FD [2]	85.14	56.81	48.54	46.70	47.72	45.82	44.57
	CAS [3]	<b>85.78</b>	55.57	50.42	49.91	<b>53.47</b>	46.46	44.23
	CIFS [4]	79.87	56.53	49.80	48.17	49.89	47.26	43.94
	FSR (Ours)	81.46	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	49.44	<b>48.34</b>	<b>46.41</b>

Table 4. Comparison of robustness (accuracy (%)) between existing methods and our method. All models are trained using AT with ResNet-18 on CIFAR-10. The best results are marked in **bold**, and more comprehensive Ensemble and AutoAttack are highlighted in grey.

Method	FGSM	PGD-20	PGD-100	C&W	Ensemble	AutoAttack
AT	56.21	48.22	46.37	47.38	45.51	44.11
+ FSR	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	<b>49.44</b>	<b>48.34</b>	<b>46.41</b>
w/o Sep	57.51	50.71	48.98	49.32	47.60	45.47
w/o Rec	57.67	50.06	48.54	49.41	47.32	44.96

Table 6. Comparison of robustness (%) of FSR applied on AT upon removing the Separation or the Recalibration stage. Model and dataset used are ResNet-18 and CIFAR-10, respectively. Best results are marked in **bold**.

[1] Madry et al., Towards deep learning models resistant to adversarial attacks. [ICLR 2018]

[2] Xie et al., Feature Denoising for Improving Adversarial Robustness. [CVPR 2019]

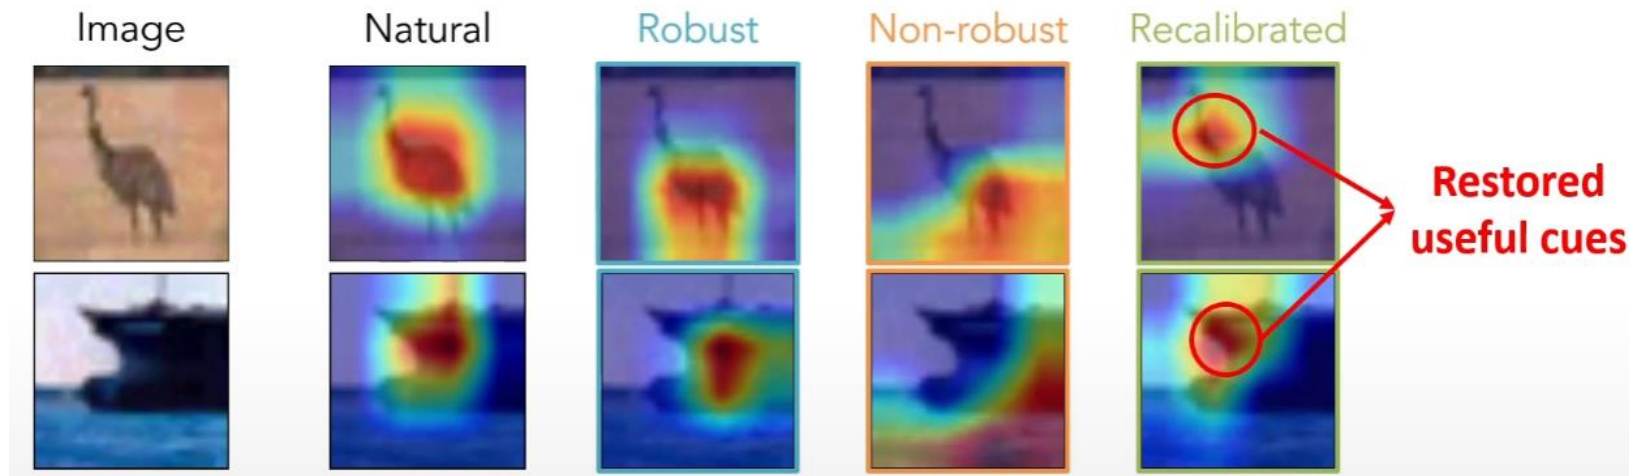
[3] Wang et al., Improving adversarial robustness via channel-wise activation suppressing. [ICLR 2021]

[4] Zhang et al., CIFS: Improving Adversarial Robustness of CNNs via Channel-wise Importance-based Feature Selection [ICML 2021]

# Experiments

## Robustness of Recalibrated Feature

Method	(a) Classification		(b) Weighted $k$ -NN	
	Ensemble	AutoAttack	5-NN	20-NN
Robust $f^+$	47.89	45.82	66.21	61.58
Non-robust $f^-$	33.11	28.39	54.69	53.89
Recalibrated $\tilde{f}^-$	46.93	44.52	66.34	65.64
$\tilde{f}^- + f^+$	48.34	46.41	70.91	65.88
$\tilde{f}$ (Ours)	48.34	46.41	70.91	65.88



# Experiments

## ■ Computational Efficiency

Method	<i>VGG16</i>		<i>ResNet-18</i>	
	# Params (M)	FLOPs (G)	# Params (M)	FLOPs (G)
Vanilla	15.25	0.6299	11.17	1.1133
+ FSR	16.52	0.6701	12.43	1.1535

Table 7. Comparison of computational costs (# params and FLOPs) on a vanilla model and a model with our FSR module.

## ■ Position of PSR module

	No attack	FGSM	PGD-20	PGD-100	C&W	Ensemble
Block1	<b>84.58</b>	56.41	48.29	46.28	46.96	44.89
Block2	83.76	56.34	48.86	47.03	47.32	45.28
Block3	82.60	56.62	50.43	49.11	47.84	46.33
Block4	81.46	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	<b>49.44</b>	<b>48.34</b>
Block3 + Block4	82.18	56.93	50.72	49.32	48.63	46.91

Table A4. Comparison of accuracy (%) as we insert our FSR module after different layers of ResNet-18.

## ■ Effects of Gumbel Softmax

	FGSM	PGD-20	PGD-100	C&W	Ensemble	AutoAttack
Binary	55.78	49.21	47.79	48.74	46.91	44.26
Gumbel	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	<b>49.44</b>	<b>48.34</b>	<b>46.41</b>

Table A6. Comparison of accuracy (%) on using mask generated by discrete binary sampling or through Gumbel softmax.

# Discussion

## ■ Limitations

- Assumption that the input images contain malicious perturbations designed to fool the model
- So, FSR module occasionally decreases the natural accuracy by a small amount

<i>ResNet-18</i>	CIFAR-10						SVHN					
	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble
AT	<b>85.02</b>	56.21	48.22	46.37	47.38	45.51	91.21	55.55	40.85	37.54	40.61	37.41
AT + FSR	81.46	<b>58.07</b>	<b>52.47</b>	<b>51.02</b>	<b>49.44</b>	<b>48.34</b>	<b>91.28</b>	<b>60.46</b>	<b>43.94</b>	<b>39.01</b>	<b>43.22</b>	<b>38.81</b>
TRADES	<b>86.31</b>	57.21	50.74	49.44	48.66	47.89	90.99	61.31	47.12	43.55	45.48	42.99
TRADES + FSR	84.49	<b>58.29</b>	<b>52.27</b>	<b>51.28</b>	<b>49.92</b>	<b>49.28</b>	<b>91.39</b>	<b>68.85</b>	<b>51.49</b>	<b>47.50</b>	<b>46.70</b>	<b>46.17</b>
MART	82.73	56.65	50.88	49.15	47.21	45.98	<b>90.50</b>	58.21	43.61	40.43	42.20	40.07
MART + FSR	<b>83.28</b>	<b>59.55</b>	<b>54.80</b>	<b>53.69</b>	<b>48.98</b>	<b>48.36</b>	89.87	<b>61.06</b>	<b>46.51</b>	<b>42.94</b>	<b>43.89</b>	<b>42.40</b>

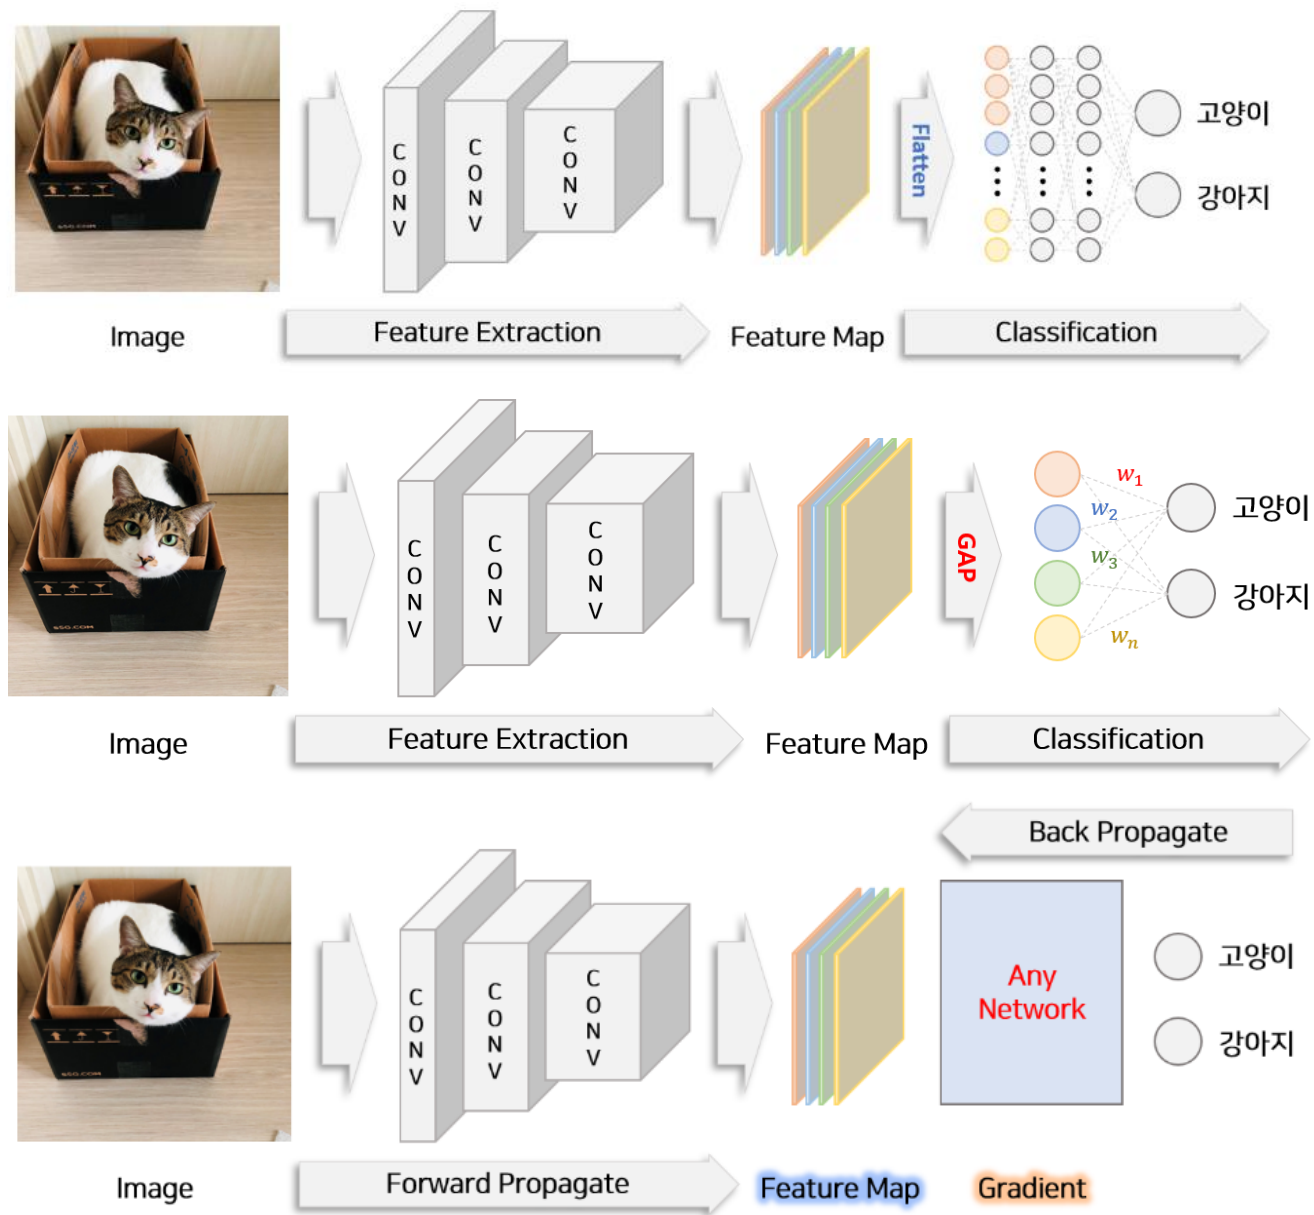
Table 1. Robustness (accuracy (%)) of adversarial training strategies (AT, TRADES, MART) with (+ FSR) and without our FSR module against diverse white-box attacks on ResNet-18. Better results are marked in **bold**.

<i>VGG16</i>	CIFAR-10						SVHN					
	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble	Natural	FGSM	PGD-20	PGD-100	C&W	Ensemble
AT	<b>80.56</b>	53.47	47.17	45.58	45.82	43.71	89.59	54.88	40.27	36.90	39.46	36.62
AT + FSR	80.06	<b>54.40</b>	<b>49.82</b>	<b>48.82</b>	<b>47.28</b>	<b>46.24</b>	<b>91.44</b>	<b>65.01</b>	<b>45.99</b>	<b>39.07</b>	<b>43.08</b>	<b>38.15</b>
TRADES	<b>82.44</b>	53.92	47.39	46.20	44.80	44.20	90.48	61.50	45.99	40.00	42.82	39.27
TRADES + FSR	80.78	<b>55.48</b>	<b>49.95</b>	<b>49.03</b>	<b>46.28</b>	<b>45.90</b>	<b>91.89</b>	<b>69.25</b>	<b>54.56</b>	<b>47.81</b>	<b>46.66</b>	<b>44.10</b>
MART	76.11	54.86	51.06	50.16	43.53	43.01	89.95	59.03	42.89	38.73	39.12	37.64
MART + FSR	<b>79.18</b>	<b>56.41</b>	<b>52.69</b>	<b>52.13</b>	<b>44.49</b>	<b>44.20</b>	<b>90.60</b>	<b>62.28</b>	<b>47.17</b>	<b>42.50</b>	<b>43.44</b>	<b>40.73</b>

Table 2. Robustness (accuracy (%)) of adversarial training strategies (AT, TRADES, MART) with (+ FSR) and without our FSR module against diverse white-box attacks on VGG16. Better results are marked in **bold**.

**Thank you!**

# Grad CAM





# Gumbel Softmax

## Gumbel-Max trick

Let  $\pi_1, \pi_2, \dots, \pi_n$  be probabilities, i.e.,  $\sum_k \pi_k = 1$

We define  $Z = \operatorname{argmax}_k \{\log \pi_k + G_k\}$  where  $G_1, \dots, G_n$  i.i.d.  $\sim \text{Gumbel}(0, 1)$

Then,  $\mathbb{P}(Z = k) = \pi_k$

Proof.

Let  $u_k = \log \pi_k + G_k$


$$\begin{aligned} \mathbb{P}(Z = k) &= \mathbb{P}(u_k \geq u_j, \forall j \neq k) \\ &= \int_{-\infty}^{\infty} \mathbb{P}(u_k \geq u_j, \forall j \neq k | u_k) \mathbb{P}(u_k) du_k \\ &= \int_{-\infty}^{\infty} \prod_{j \neq k} \mathbb{P}(u_k \geq u_j | u_k) \mathbb{P}(u_k) du_k \\ &= \int_{-\infty}^{\infty} \prod_{j \neq k} e^{-e^{-u_k + \log \pi_j}} e^{-(u_k - \log \pi_k + e^{-(u_k - \log \pi_k)})} du_k \\ &= \int_{-\infty}^{\infty} e^{-\sum_{j \neq k} \pi_j e^{-u_k}} \pi_k e^{-(u_k + \pi_k e^{-u_k})} du_k \\ &= \pi_k \int_{-\infty}^{\infty} e^{-u_k - (\pi_k + \sum_{j \neq k} \pi_j) e^{-u_k}} du_k = \pi_k \end{aligned}$$

Gumbel(0, 1)

PDF  $f(x) = e^{-x} e^{-e^{-x}}$

CDF  $F(x) = e^{-e^{-x}}$

## Gumbel-Max Trick

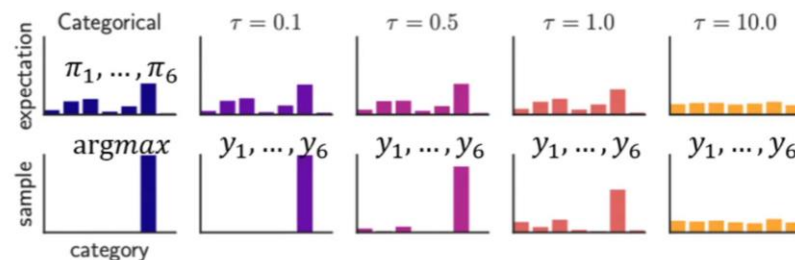


(a) Discrete( $\alpha$ )


## Gumbel-Softmax trick

$$\operatorname{argmax}_k \{\log \pi_k + G_k\} \rightarrow y_i = \frac{e^{(\log \pi_k + G_k)/\tau}}{\sum_{k=1}^n e^{(\log \pi_k + G_k)/\tau}}$$

for  $i = 1, \dots, n$  where  $\tau > 0$  is softmax temperature



## Gumbel-Softmax Trick



(b) Concrete( $\alpha, \lambda$ )